

Self-Distilled Agentic Reinforcement Learning: a Literature-Situated Review

Paul Leyva

Study repo: `ai-research-studies/self-distilled-agentic-reinforcement-learning`

May 16, 2026

Notation

A compact symbol table is reproduced from `02-math-deep-dive.md`; full glossary is in that file.

Symbol	Meaning
$\pi_\theta(\cdot \mid s_t)$	Student policy on student context s_t .
$\pi_T(\cdot \mid s_t^+)$	Teacher branch: same parameters θ , privileged context s_t^+ .
y_t	Token sampled by the student at position t .
Δ_t	Teacher-student log-prob gap, $\log \pi_T(y_t \mid s_t^+) - \log \pi_\theta(y_t \mid s_t)$.
h_t	Student entropy, $-\sum_v \pi_\theta(v \mid s_t) \log \pi_\theta(v \mid s_t)$.
$g_t \in [0, 1]$	Per-token sigmoid gate (stop-gradient).
β	Gate sharpness (overloaded with the GRPO KL coefficient).
λ_{SDAR}	Auxiliary loss weight on the SDAR term.

Background: Multi-Turn Agentic Post-Training

Modern post-training of LLM agents is dominated by reinforcement learning with verifiable rewards. Group Relative Policy Optimization (GRPO), introduced by [Shao et al. \[2024\]](#) and scaled by [DeepSeek-AI \[2025\]](#), is now the workhorse: it strips the value head from PPO [[Schulman et al., 2017](#)] and replaces the GAE baseline with a group-mean baseline computed across G rollouts of the same prompt. The trajectory-level reward is then broadcast token-wise via a clipped importance ratio. This pipeline is conceptually a descendant of RLHF [[Christiano et al., 2017](#), [Ouyang et al., 2022](#)], with the human-preference reward model swapped for a programmatic verifier.

The structural weakness of GRPO in agentic settings is supervisory sparsity. A 20-turn trajectory in ALFWorld [[Shridhar et al., 2021](#)] or WebShop [[Yao et al., 2022](#)] compresses thousands of token-level decisions into a single scalar reward. *On-Policy Distillation* [[Ye et al., 2026](#), [Yang et al., 2026](#)] and its self-distilled variant *OPSD* [[Zhao et al., 2026](#)] address this by adding a per-token auxiliary loss against a teacher branch. In the self-distilled regime the teacher is the *same* policy conditioned on a privileged context $s_t^+ = s_t \cup$ (retrieved skill text) available only at training time — a free dense signal. SDAR [[Lu et al., 2026](#)] is the agentic extension of this idea, plus a sigmoid gate to handle the multi-turn failure modes documented below.

Method: SDAR

The full objective is

$$\mathcal{L}(\theta) = \mathcal{L}_{\text{GRPO}}(\theta) + \lambda_{\text{SDAR}} \cdot \mathcal{L}_{\text{SDAR}}(\theta), \quad (1)$$

with the per-token SDAR loss

$$\ell_t^{\text{SDAR}} = g_t \cdot (\log \pi_T(y_t | s_t^+) - \log \pi_\theta(y_t | s_t)) = g_t \cdot \Delta_t, \quad (2)$$

aggregated by masked-token average, $\mathcal{L}_{\text{SDAR}} = (\sum_t m_t g_t \Delta_t) / (\sum_t m_t)$. The gap Δ_t is exactly the negative single-sample estimator of the per-token reverse KL between teacher and student, evaluated at the student-sampled token — avoiding a full $|\mathcal{V}|$ -wide softmax sum.

Lu et al. [2026] benchmark three gating strategies:

$$\text{entropy: } g_t = \sigma(\beta h_t), \quad (3)$$

$$\text{gap: } g_t = \sigma(\beta \Delta_t), \quad (4)$$

$$\text{soft-OR: } g_t = \sigma(\beta [1 - (1 - h_t)(1 - \Delta_t)]). \quad (5)$$

Gap gating is the conceptual centerpiece: it operationalises the asymmetric-trust observation. When $\Delta_t > 0$ the privileged context is endorsing the student’s sampled token — high-value distillation signal — and the gate opens. When $\Delta_t < 0$ the disagreement is ambiguous (could be skill-retrieval failure rather than a true student error), and the sigmoid attenuates the update. The gate is detached, so the gradient flows only through the student log-prob:

$$\nabla_\theta \mathcal{L}_{\text{SDAR}} = -\text{Agg}[g_t \cdot \nabla_\theta \log \pi_\theta(y_t | s_t)], \quad (6)$$

which is structurally a gated REINFORCE-style update with g_t as a per-token signed weight.

Results

We summarise Table 1 of Lu et al. [2026] on Qwen 2.5 [Qwen Team, 2024].

Method (Qwen 2.5-3B-Instruct)	ALFWorld $\uparrow$	Search-QA $\uparrow$	WebShop $\uparrow$
GRPO baseline	75.0	36.4	79.8
GRPO + OPSD (ungated)	—	—	—
Skill-GRPO	—	—	—
Skill-SD	—	—	—
RLSD	—	—	—
SDAR (ours)	84.4	43.4	85.0
Numbers reproduced from 02-math-deep-dive.md §8; dashes indicate baseline numbers not transcribed into the study notes.			

The headline gains are +9.4 pp on ALFWorld average success, +7.0 pp on Search-QA [Dunn et al., 2017], and +5.2 pp on WebShop Score for the 3B model. Qwen 2.5-7B-Instruct shows comparable gains. The qualitative companion result is Figure 2 of Lu et al. [2026]: ungated GRPO+OPSD exhibits catastrophic KL-divergence growth across training, while gated SDAR keeps the per-turn KL bounded — direct empirical support for the multi-turn-instability story.

Position in the Literature

RL backbone. SDAR’s GRPO term is taken essentially verbatim from Shao et al. [2024]. The PPO surrogate [Schulman et al., 2017] is therefore the deepest underlying optimisation primitive — with the value head removed and the GAE baseline replaced by a group-mean baseline. SDAR contributes *nothing* to the RL backbone itself; its novelty is the auxiliary objective.

Distillation lineage. The auxiliary term descends from On-Policy Distillation [Ye et al., 2026, Yang et al., 2026], which formalised teacher guidance on student-sampled trajectories, and from On-Policy Self-Distillation [Zhao et al., 2026], which collapses teacher and student into a single set of parameters with a privileged-context teacher branch. SDAR’s contribution to this lineage is two-fold: (i) the agentic, multi-turn observation that vanilla OPSD diverges, and (ii) the gating remedy.

Curriculum vs gating. A parallel response to multi-turn OPSD instability is the turn-level curriculum of TCOB [Wang et al., 2026], which schedules teacher influence by turn index. SDAR’s gating is the “smooth” alternative — the gate is computed per token from the realised teacher–student gap rather than scheduled exogenously by turn. The two designs are not mutually exclusive; combining them is an obvious follow-up.

RLHF foundations. Conceptually the gated distillation term plays a role analogous to the KL-to-reference regulariser in InstructGPT [Ouyang et al., 2022] and the broader RLHF programme [Christiano et al., 2017]: both prevent the policy from drifting too far from a reference distribution. SDAR replaces the static SFT reference with a dynamic privileged-context teacher and gates the constraint per-token rather than treating it as a fixed Lagrangian.

Sibling agentic-RL work. Hegde and NovaSky-AI Contributors [2026] (the “Reinforcing Recursive Language Models” alphaXiv blog from the NovaSky-AI / SkyRL group) attacks the same supervisory-sparsity problem from a different angle: a shared policy is rolled out across a *recursive call tree*, and child rollouts inherit advantages from their parent. Credit assignment is by tree topology rather than by per-token reliability gating. The two designs are complementary; § Discussion sketches a hybrid.

Benchmarks. The agentic suite is standard: ALFWorld [Shridhar et al., 2021], WebShop [Yao et al., 2022], and Search-QA-style tasks [Dunn et al., 2017]. None of these were designed to stress 50+-turn horizons, which is one reason the depth-scaling question (below) remains open.

Discussion

What the experiments support. Three claims are well-supported by Table 1 and Figure 2 of Lu et al. [2026]: (a) ungated OPSD on top of GRPO destabilises multi-turn training; (b) any of the three proposed gates restores stability; (c) gated SDAR strictly improves over the GRPO baseline on every reported aggregate metric for both 3B and 7B backbones. The KL-bounded behaviour shown in Figure 2 is direct evidence for the failure-mode analysis in § 1 of the paper.

What is hand-waved. The bias–variance argument for gating is informal in the paper — there is no formal theorem that gap-gated OPSD is a lower-variance estimator than ungated OPSD, only an empirical demonstration. Sensitivity to the two new hyperparameters (β and λ_{SDAR}) is not ablated in the main text. The interpretation of Δ_t as a “trust” signal also presupposes that teacher log-prob is a reliable proxy for skill-retrieval quality — a confidently wrong teacher (high Δ_t , low actual quality) would be heavily distilled into the student. This is plausible but not directly tested; the slight under-performance of SDAR vs GRPO+OPSD on some Search-QA sub-tasks is consistent with this critique.

Three future directions. (1) *Depth scaling*: re-run on a benchmark with 50–100 effective turns and check whether the per-turn KL stays bounded. (2) *Adaptive gates*: learn β jointly with θ , or replace the hand-designed gate with $g_\phi(s_t, \Delta_t, h_t)$. (3) *Recursive-call extension*: lift gating into the recursive setting of [Hegde and NovaSky-AI Contributors \[2026\]](#) — gate child rollouts by whether the parent endorses them, combining SDAR-style per-token gating with RLM-style tree-level credit assignment.

See Also

This paper sits inside a broader study chain. The mathematical prerequisites are derived from scratch in [Chapter 25 \(causal LM as MDP\)](#), [Chapter 27 \(tiny GPT pre-training\)](#), [Chapter 28 \(SFT, RLHF, DPO\)](#), [Chapter 30 \(max-entropy RL\)](#), and [Chapter 31 \(PG / GRPO derivation\)](#). The sibling agentic-RL study lives at [pleyva2004/reinforcing-recursive-language-models](#) — it tackles the same supervisory-sparsity problem with a different mechanism (shared-policy + advantage inheritance across a recursive call tree, vs SDAR’s teacher branch + gated per-token distillation), and a hybrid is sketched above.

Reproduction

A locally-verified empirical reproduction of three paper claims (SDAR vs GRPO reward / KL on the toy MDP; entropy-gating KL; ungated GRPO+OPSD collapse) and two measurements of proposed improvements (λ_{SDAR} sensitivity sweep; adaptive-quantile gating tradeoff) is recorded at [findings.md](#) in the study repository. The findings table includes provenance (machine, numpy version, seed) and per-finding reproduction commands.

References

- Paul Christiano, Jan Leike, Tom B. Brown, Miljan Martic, Shane Legg, and Dario Amodei. Deep reinforcement learning from human preferences. *Advances in Neural Information Processing Systems*, 30, 2017. URL <https://arxiv.org/abs/1706.03741>.
- DeepSeek-AI. DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025. URL <https://arxiv.org/abs/2501.12948>.
- Matthew Dunn, Levent Sagun, Mike Higgins, V. Ugur Guney, Volkan Cirik, and Kyunghyun Cho. SearchQA: A new Q&A dataset augmented with context from a search engine. *arXiv preprint arXiv:1704.05179*, 2017. URL <https://arxiv.org/abs/1704.05179>.

- Niranjan Hegde and NovaSky-AI Contributors. Reinforcing recursive language models. alphaXiv blog / SkyRL, 2026. URL <https://github.com/pleyva2004/reinforcing-recursive-language-models>. TODO: verify exact arxiv id; sibling agentic-RL work, shared-policy + advantage inheritance across recursive call tree.
- Zhengxi Lu, Zhiyuan Yao, Zhuowen Han, Zi-Han Wang, Jinyang Wu, et al. Self-distilled agentic reinforcement learning. *arXiv preprint arXiv:2605.15155*, 2026. URL <https://arxiv.org/abs/2605.15155>. ZJU-REAL Lab and Meituan. Code: <https://github.com/ZJU-REAL/SDAR>.
- Long Ouyang, Jeff Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 35, 2022. URL <https://arxiv.org/abs/2203.02155>.
- Qwen Team. Qwen2.5: A party of foundation models. Technical Report, 2024. URL <https://arxiv.org/abs/2412.15115>.
- John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017. URL <https://arxiv.org/abs/1707.06347>.
- Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Mingchuan Zhang, Y. K. Li, Y. Wu, and Daya Guo. DeepSeekMath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024. URL <https://arxiv.org/abs/2402.03300>. Introduces Group Relative Policy Optimization (GRPO).
- Mohit Shridhar, Xingdi Yuan, Marc-Alexandre Côté, Yonatan Bisk, Adam Trischler, and Matthew Hausknecht. ALFWorld: Aligning text and embodied environments for interactive learning. In *International Conference on Learning Representations*, 2021. URL <https://arxiv.org/abs/2010.03768>. Preprint posted 2020.
- Jiajun Wang et al. TCOd: Turn-curriculum on-policy distillation for multi-turn agents. *arXiv preprint arXiv:2604.XXXXX*, 2026. TODO: verify; cited as Wang et al. 2026b in SDAR. Curriculum-learning alternative for multi-turn OPSD instability.
- Yifan Yang et al. On-policy distillation revisited. *arXiv preprint arXiv:2604.XXXXX*, 2026. TODO: verify; cited as Yang et al. 2026b in SDAR.
- Shunyu Yao, Howard Chen, John Yang, and Karthik Narasimhan. WebShop: Towards scalable real-world web interaction with grounded language agents. In *Advances in Neural Information Processing Systems*, volume 35, 2022. URL <https://arxiv.org/abs/2207.01206>.
- Qinyuan Ye et al. On-policy distillation for language models. *arXiv preprint arXiv:2603.XXXXX*, 2026. TODO: verify author list and arxiv id; cited as Ye et al. 2026 in SDAR.
- Hao Zhao et al. On-policy self-distillation for language-model post-training. *arXiv preprint arXiv:2602.XXXXX*, 2026. TODO: verify; cited as Zhao et al. 2026 (OPSD) in SDAR.