

Improvements and Extensions to Self-Distilled Agentic Reinforcement Learning (SDAR)

Paul Leyva

ai-research-studies/self-distilled-agentic-reinforcement-learning

2026-05-16

This document proposes four concrete extensions to SDAR [Lu et al., 2026]. Each section follows the same template: (a) statement, (b) prior-work anchoring, (c) what would test it, (d) prototype pointer. Notation matches the math deep dive (02-math-deep-dive.md); in particular, $\Delta_t = \log \pi_T(y_t | s_t^+) - \log \pi_\theta(y_t | s_t)$ is the teacher–student gap, $g_t = \sigma(\beta \text{score}_t)$ is the detached gate, and the SDAR objective is $\mathcal{L}(\theta) = \mathcal{L}_{\text{GRPO}}(\theta) + \lambda_{\text{SDAR}} \cdot \mathcal{L}_{\text{SDAR}}(\theta)$ with $\mathcal{L}_{\text{SDAR}} = \text{Agg}(g_t \Delta_t)$.

1. Mathematical Improvements: a bias bound for gap gating

Statement. Section 7 of the math deep dive sketches that gap-gating reduces bias relative to ungated OPSD, but no closed form is given. We propose the following.

Assumption 1 (Asymmetric reliability of the gap distribution). *Let $p(\Delta)$ denote the marginal distribution of Δ_t over training tokens (under the joint trajectory distribution induced by π_θ and the environment). There exists a threshold $\tau \leq 0$ such that the mass $\int_{-\infty}^{\tau} p(\Delta) d\Delta$ is dominated by skill-retrieval failures — the privileged context s_t^+ contains text that misled the teacher branch, so $\Delta < \tau$ is unreliable distillation signal. Mass with $\Delta \geq \tau$ corresponds to reliable endorsement (or near-zero disagreement).*

This is the formalisation of paper §1’s three failure modes (skill quality, skill utilisation, multi-turn drift), all of which manifest as *negative-tail* mass in $p(\Delta)$.

Proposition 1 (Bias bound for the gated estimator). *Let $\hat{g}_{\text{ungated}}(\theta) = \mathbb{E}_{\Delta \sim p}[\Delta]$ be the per-token mean of the ungated SDAR contribution and let $\hat{g}_{\text{gated}}(\theta) = \mathbb{E}_{\Delta \sim p}[g(\Delta)\Delta]$ with $g(\Delta) = \sigma(\beta\Delta)$. Define the ideal estimator $\hat{g}_{\text{true}}(\theta) = \int_{\tau}^{\infty} \Delta p(\Delta) d\Delta$ (the contribution restricted to reliable mass). Then*

$$|\hat{g}_{\text{gated}}(\theta) - \hat{g}_{\text{true}}(\theta)| \leq \underbrace{\int_{-\infty}^{\tau} (1-g(\Delta)) |\Delta| p(\Delta) d\Delta}_{\text{negative-tail leakage}} + \underbrace{\int_{\tau}^{\infty} (1-g(\Delta)) \Delta p(\Delta) d\Delta}_{\text{reliable-tail attenuation}}.$$

Both terms vanish as $\beta \rightarrow \infty$ (hard step at zero), with rate $O(\exp(-\beta|\tau|))$ for the first and $O(\beta^{-1})$ for the second.

Prior work. The decomposition mirrors the baseline-doesn’t-bias proof of GRPO [Shao et al., 2024] and the chain Ch. 31 derivation, where a state-dependent baseline is shown to be zero-mean against the policy gradient. The gate g_t is detached, so the same conditional-expectation argument applies: gating shifts the *target* but does not introduce a new dependence on θ . The asymmetric-tail framing is inspired by importance-sampling truncation analyses common in PPO [Schulman et al., 2017] and InstructGPT-style RLHF [Ouyang et al., 2022].

What would test it. The full statement and proof appear in `proofs/gating-bias-bound.tex`. Empirically, the toy bandit in `sandbox/toy_sdar.py` (and the comparison in `improvements/adaptive-gate.py`) lets us measure the empirical bias of the gated vs. ungated estimator under a synthetic $p(\Delta)$ with controllable negative-tail mass.

Prototype pointer. PROOF: `proofs/gating-bias-bound.tex` (stub). Numerical sanity check: `improvements/adaptive-gate.py` already computes per-token gate-fire rates; extending to bias estimates is ~ 10 lines.

2. Code / Implementation Improvements: adaptive-quantile gating

Statement. The paper uses a fixed sigmoid $g_t = \sigma(\beta \Delta_t)$ centred at $\Delta_t = 0$. We propose

$$g_t = \sigma(\beta(\Delta_t - \hat{q}_\tau)), \quad \hat{q}_\tau \leftarrow (1 - \eta) \hat{q}_\tau + \eta \Delta_t^{\text{thresh}},$$

where $\hat{q}_\tau$ is an exponentially-moving estimate of the τ -quantile (e.g. $\tau = 0.3$) of recent gaps. Rationale: the gap distribution drifts over training as the student policy improves; the *absolute* criterion $\Delta_t > 0$ stops being meaningful after a few hundred steps because typical gaps shift toward zero. A relative criterion “fire on the top $1 - \tau$ fraction” stays well-calibrated.

Prior work. Quantile-based truncation has a long history in robust statistics; the closest deep-RL precedent is V-TRACE-style importance-clipping in IMPALA, and the running-quantile baselines in Wang et al. [2026]. The detached-gate structure is preserved, so the convergence properties argued in Prop. 1 still apply.

What would test it. Three-way comparison on the toy MDP: GRPO baseline, fixed-threshold SDAR ($\sigma(2\Delta_t)$), and adaptive-quantile SDAR. Metric: final reward subject to a per-turn KL budget identical to the paper’s stability criterion.

Prototype pointer. MEASUREMENT: `improvements/adaptive-gate.py`. Returns `{fixed_threshold_rewa, adaptive_kl, gate_fire_rate_fixed, gate_fire_rate_adaptive}`.

3. Experimental Extensions: λ_{SDAR} sensitivity sweep

Statement. The SDAR ablation in the paper fixes λ_{SDAR} at one value and varies the gating strategy. But λ_{SDAR} is itself a critical hyperparameter: $\lambda_{\text{SDAR}} = 0$ recovers GRPO and gives no distillation; $\lambda_{\text{SDAR}} \rightarrow \infty$ overwhelms the RL signal and the policy collapses to a maximum-likelihood estimate of the teacher. Somewhere between is an optimum. We propose a 5-point

sweep $\lambda_{\text{SDAR}} \in \{0, 0.1, 0.3, 1, 3\}$ at fixed β , reporting both final reward and final per-turn KL, and identifying the constrained optimum

$$\lambda^* = \arg \max_{\lambda} \text{reward}(\lambda) \quad \text{s.t. } \text{kl}(\lambda) \leq 2 \cdot \text{kl}(0).$$

Prior work. Standard auxiliary-loss-weight sweeps appear in DeepSeek-AI [2025] and Zhao et al. [2026]; both observe an inverted-U in reward as the auxiliary weight grows. Our sweep is the agentic-RL specialisation of this exercise.

What would test it. Five 100-step training runs on the toy MDP, one per λ value, with $\beta = 2$ held constant. Constrained arg max over the resulting (reward, KL) Pareto trace.

Prototype pointer. MEASUREMENT: `improvements/lambda-sweep.py`. Returns the full `{sweep, optimal_lambda}` dictionary.

4. Theoretical / Conceptual Connections: SDAR as constrained RL

Statement. Frame the SDAR objective as a *constrained* GRPO problem with a per-token trust constraint:

$$\max_{\theta} J_{\text{GRPO}}(\theta) \quad \text{s.t. } \mathbb{E}_t[\mathbf{1}[t \in \mathcal{T}] D_{\text{KL}}(\pi_{\theta}(\cdot | s_t) \| \pi_T(\cdot | s_t^+))] \leq \delta,$$

where $\mathcal{T}$ is the set of *trusted* tokens (those for which the privileged context endorses the student trajectory). The Lagrangian is

$$\mathcal{L}(\theta, \mu) = J_{\text{GRPO}}(\theta) - \mu \cdot \left(\mathbb{E}_t[\mathbf{1}[t \in \mathcal{T}] \cdot D_{\text{KL}}^{(t)}] - \delta \right).$$

SDAR’s gated loss has *exactly* this structure with the global multiplier μ replaced by a state-dependent per-token multiplier $\lambda_{\text{SDAR}} \cdot g_t$. In other words, SDAR is a *soft, primal-only* approximation to a per-token trust-constrained RL problem, where the gate g_t is an indicator surrogate for membership in $\mathcal{T}$ and λ_{SDAR} is a fixed multiplier that plays the role of the dual variable.

Prior work. Constrained RL with primal-dual saddle-point solvers traces back to the option/sub-goal formalism of ?. The chain Ch. 30 max-entropy-RL derivation gives the same Lagrangian form for the entropy-regularised case. The modern primal-only variant (treating μ as a hyperparameter rather than a learned dual variable) is the same shortcut taken by PPO’s KL-penalty schedule [Schulman et al., 2017].

What would test it. This is a translation, not a new theorem. The deliverable is a short proof sketch in `proofs/sdar-as-constrained-rl.tex` that (i) writes the SDAR gradient as the gradient of the Lagrangian above with $\mu = \lambda_{\text{SDAR}}$ and the indicator replaced by g_t , and (ii) cites the saddle-point convergence conditions from constrained-RL literature for which SDAR’s fixed multiplier is sound (essentially: when the constraint is slack at the optimum, primal-only with a sufficiently large μ works).

Prototype pointer. PROOF: `proofs/sdar-as-constrained-rl.tex` (stub).

Discussion

The four improvements compose. The mathematical bias bound (§1) is the *justification* for adaptive-quantile gating (§2): the bound’s negative-tail leakage term shrinks faster when the gate centre tracks the empirical τ -quantile rather than sitting at zero. The λ_{SDAR} sweep (§3), once run jointly with the adaptive gate, would reveal a *joint* optimum (λ^*, τ^*) , plausibly different from the single-variable optimum. And the constrained-RL framing (§4) gives the λ_{SDAR} sweep its theoretical reading: the optimal λ^* is an estimate of the dual variable that would solve the constrained problem at the desired KL budget. Together they upgrade SDAR from a heuristic auxiliary loss to a principled primal approximation of a per-token trust-constrained agentic-RL problem.

References

- DeepSeek-AI. DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025. URL <https://arxiv.org/abs/2501.12948>.
- Zhengxi Lu, Zhiyuan Yao, Zhuowen Han, Zi-Han Wang, Jinyang Wu, et al. Self-distilled agentic reinforcement learning. *arXiv preprint arXiv:2605.15155*, 2026. URL <https://arxiv.org/abs/2605.15155>. ZJU-REAL Lab and Meituan. Code: <https://github.com/ZJU-REAL/SDAR>.
- Long Ouyang, Jeff Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 35, 2022. URL <https://arxiv.org/abs/2203.02155>.
- John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017. URL <https://arxiv.org/abs/1707.06347>.
- Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Mingchuan Zhang, Y. K. Li, Y. Wu, and Daya Guo. DeepSeekMath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024. URL <https://arxiv.org/abs/2402.03300>. Introduces Group Relative Policy Optimization (GRPO).
- Jiajun Wang et al. TCOd: Turn-curriculum on-policy distillation for multi-turn agents. *arXiv preprint arXiv:2604.XXXXX*, 2026. TODO: verify; cited as Wang et al. 2026b in SDAR. Curriculum-learning alternative for multi-turn OPSD instability.
- Hao Zhao et al. On-policy self-distillation for language-model post-training. *arXiv preprint arXiv:2602.XXXXX*, 2026. TODO: verify; cited as Zhao et al. 2026 (OPSD) in SDAR.